

Part II

Process Management


Chapter 4: Threads

What Is a Thread?

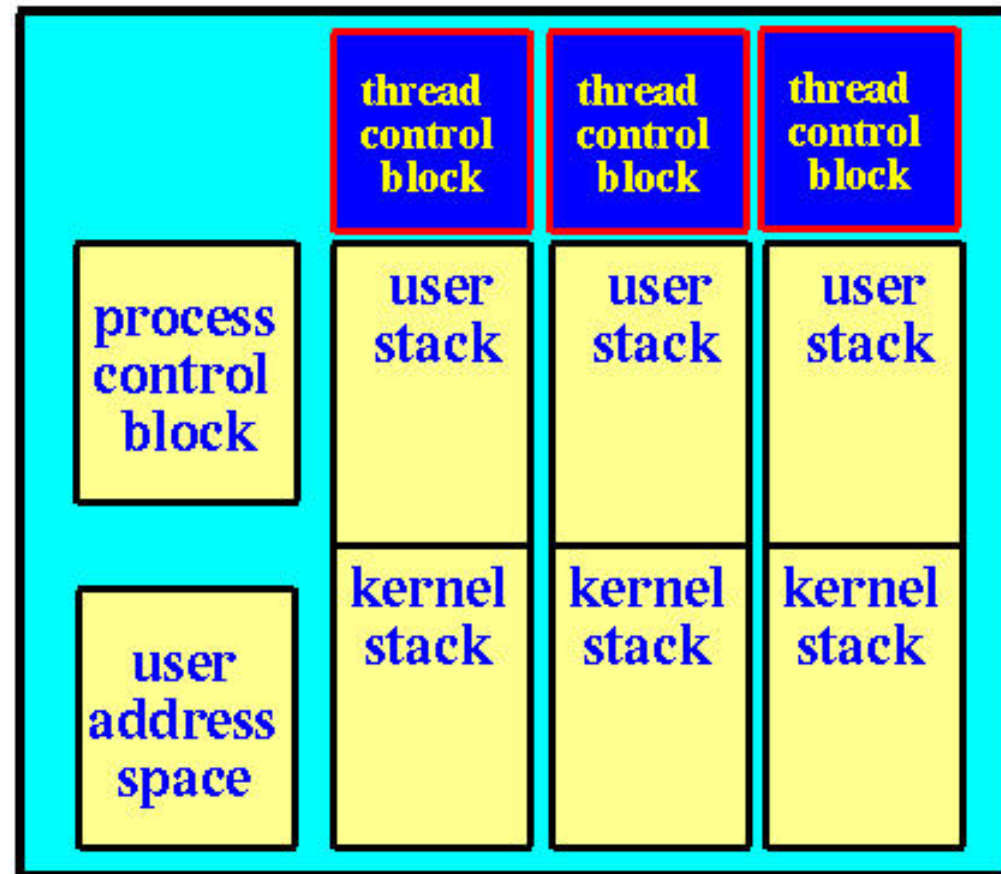
- A *thread*, also known as *lightweight process* (LWP), is a basic unit of CPU execution.
- A thread has a **thread ID**, a **program counter**, a **register set**, and a **stack**. Thus, it is similar to a process.
- However, a thread *shares* with other threads in the *same* process its code section, data section, and other OS resources (*e.g.*, files and signals).
- A process, or **heavyweight** process, has a *single* thread of control.

Single Threaded and Multithreaded Process

Single-threaded process



Multithreaded Process



Benefits of Using Threads

- **Responsiveness:** Other parts (*i.e.*, threads) of a program may still be running even if one part (*e.g.*, a thread) is blocked.
- **Resource Sharing:** Threads of a process, by default, share many system resources (*e.g.*, files and memory).
- **Economy:** Creating and terminating processes, allocating memory and resources, and context switching processes are very time consuming.
- **Utilization of Multiprocessor Architecture:** Multiple CPUs can run multiple threads of the same process. No program change is necessary.

User and Kernel Threads: 1/3

□ *User Threads:*


- ❖ User threads are supported at the **user level**. The kernel **is not aware** of user threads.
- ❖ A library provides all support for thread creation, termination, joining, and scheduling.
- ❖ There is no kernel intervention, and, hence, user threads are usually **more efficient**.
- ❖ Unfortunately, since the kernel only recognizes the containing process (of the threads), *if one thread is blocked, all threads of the same process are also blocked* because the containing process is blocked.

User and Kernel Threads: 2/3

□ *Kernel threads:*

- ❖ Kernel threads are supported by the kernel. The kernel does thread creation, termination, joining, and scheduling in kernel space.
- ❖ Kernel threads are usually **slower** than user threads due to system overhead.
- ❖ However, *blocking one thread will not cause other threads of the same process to block*. The kernel simply runs other threads.
- ❖ In a multiprocessor environment, the kernel can schedule threads on different processors.

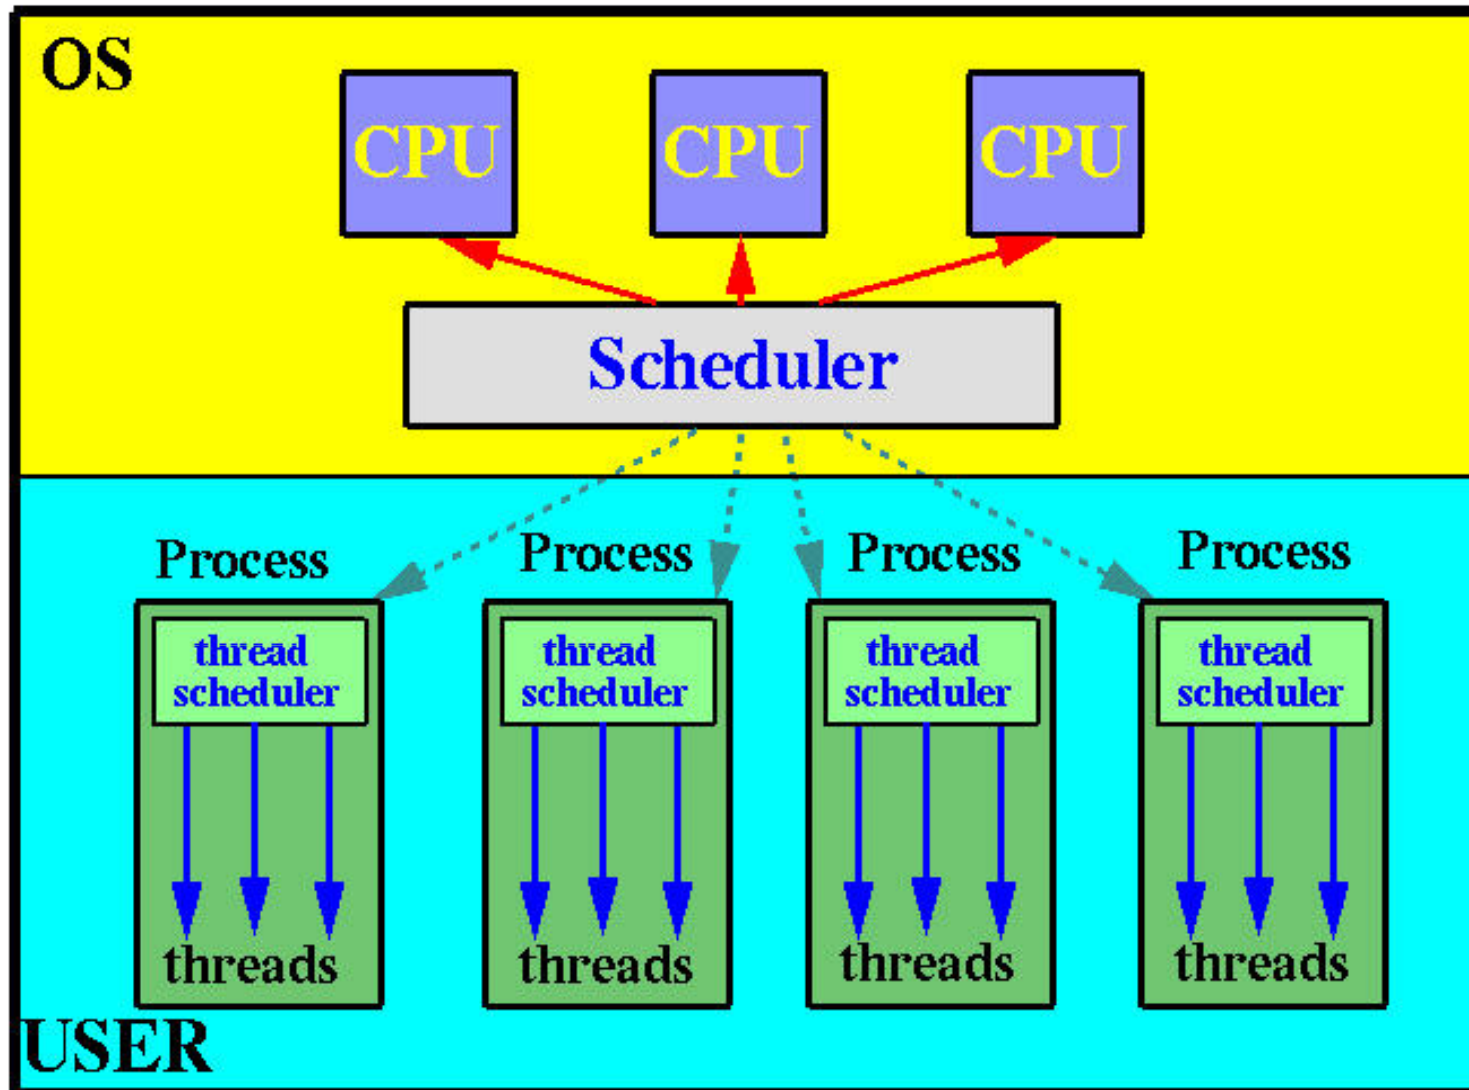
User and Kernel Threads: 3/3



Multithreading Models

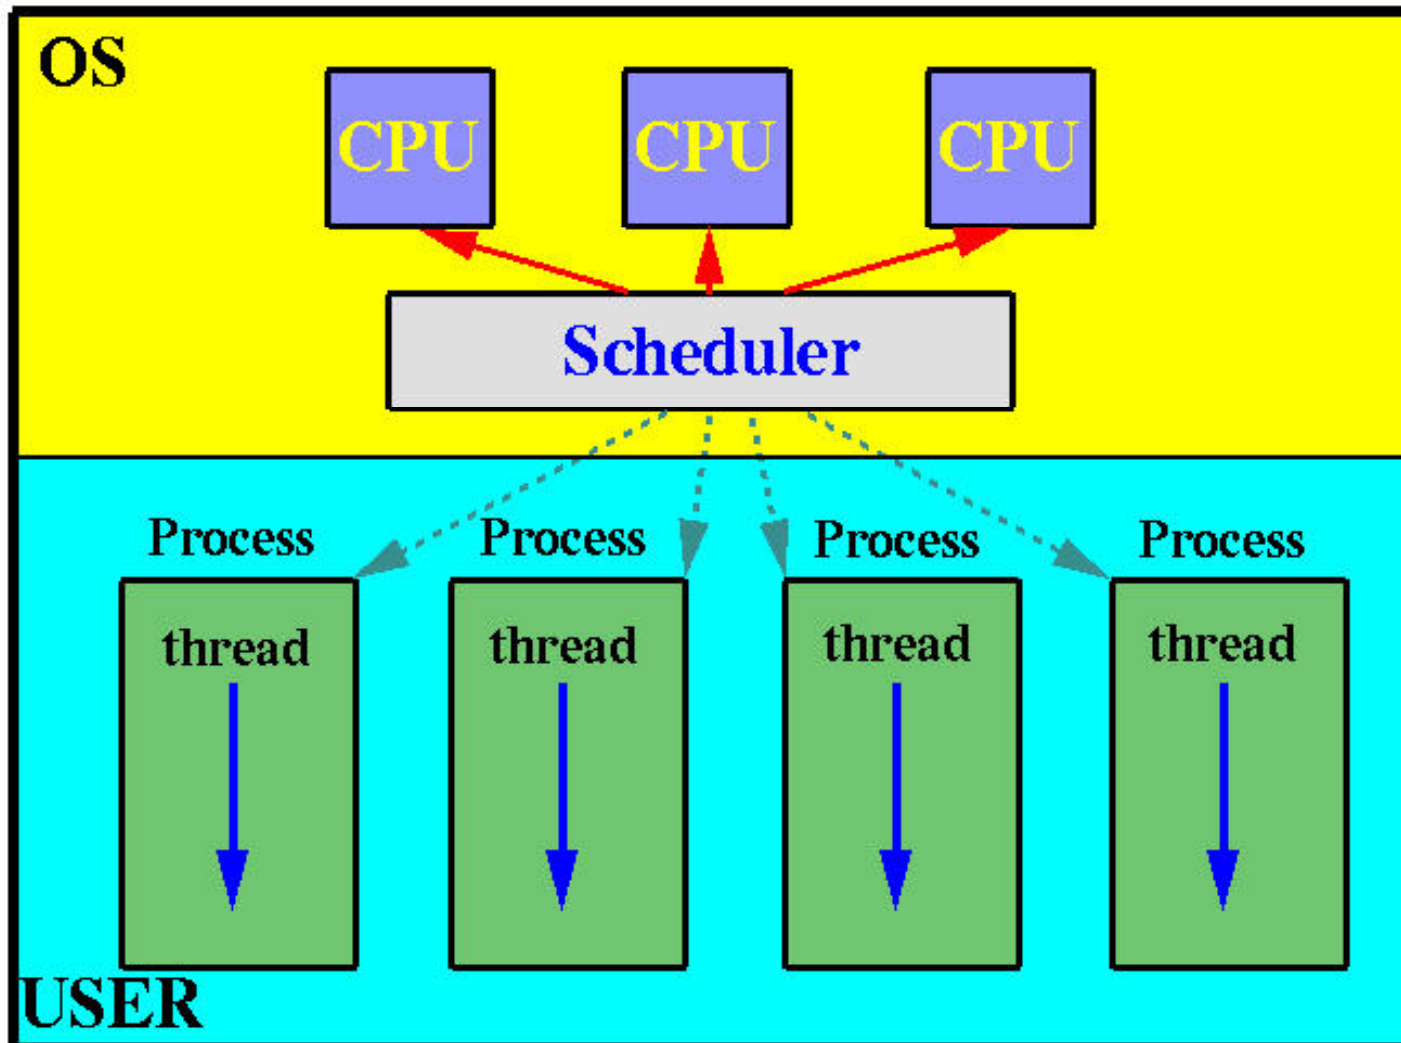
- ❑ Different systems support threads in different ways. Here are three commonly seen thread models:
 - ❖ ***Many-to-One Model***: One kernel thread (or process) has multiple user threads. Thus, this is a user thread model.
 - ❖ ***One-to-One Model***: One user thread maps to one kernel thread (*e.g.*, old Unix/Linux and Windows systems).
 - ❖ ***Many-to-Many Model***: Multiple user threads maps to a number of kernel threads.

Many-to-One Model

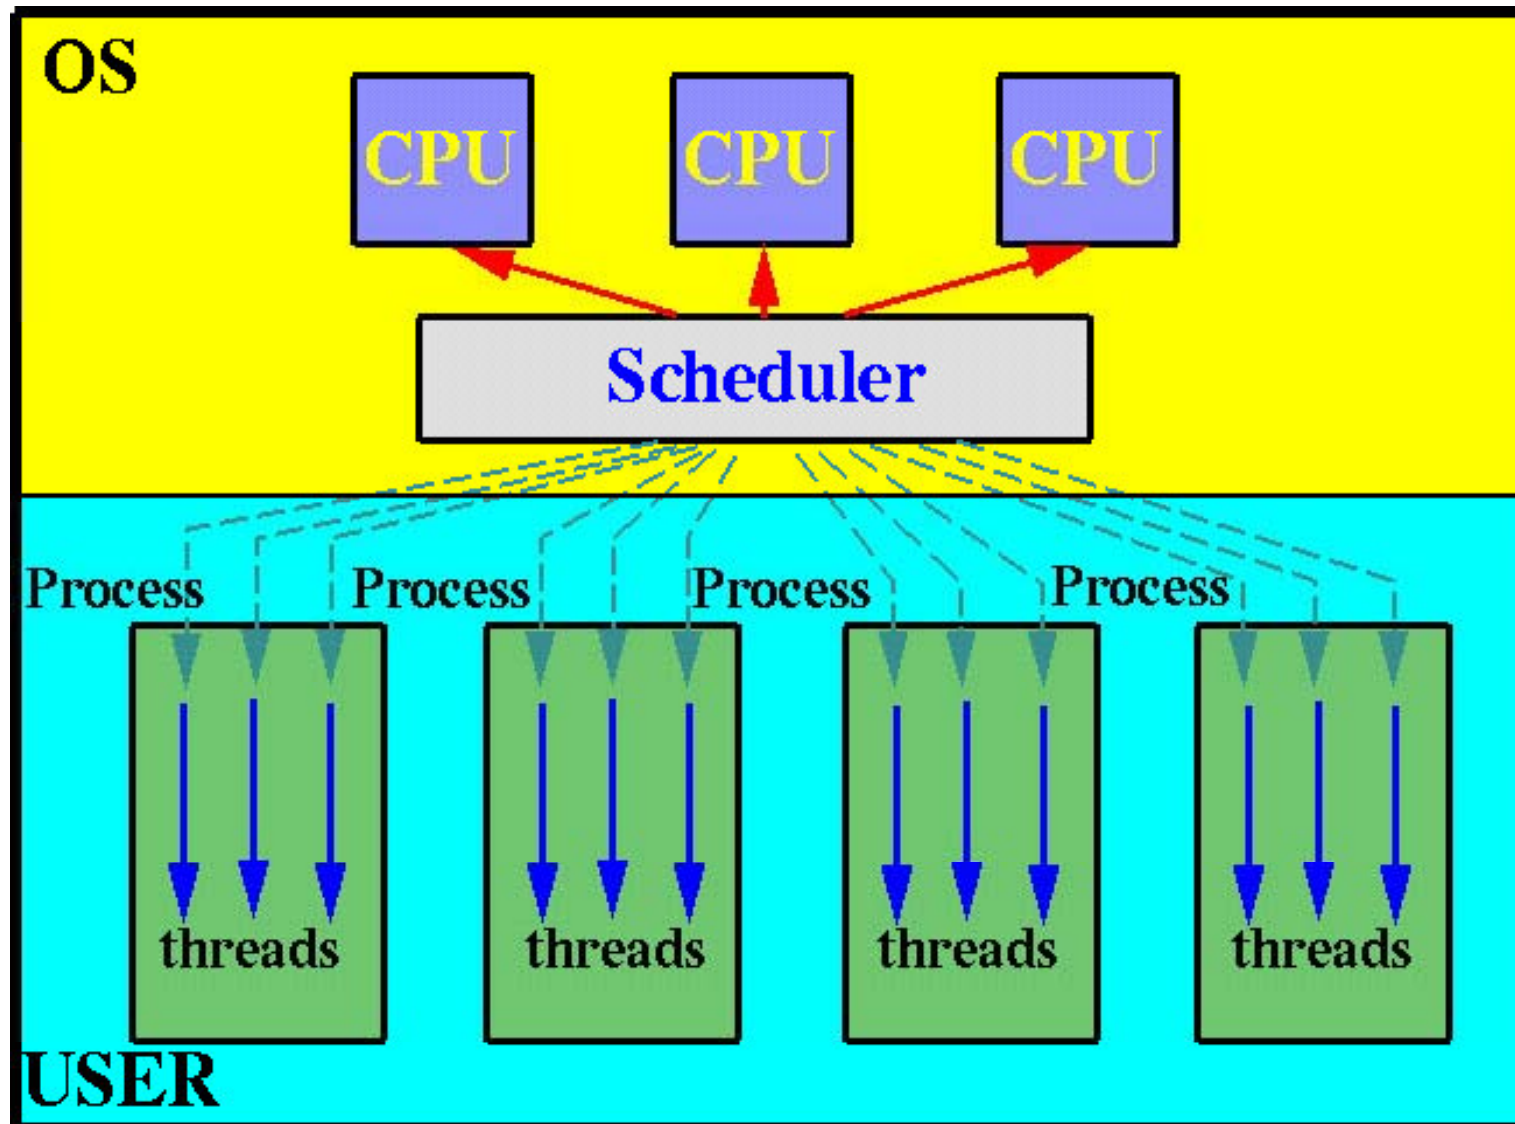


One-to-One Model: 1/2

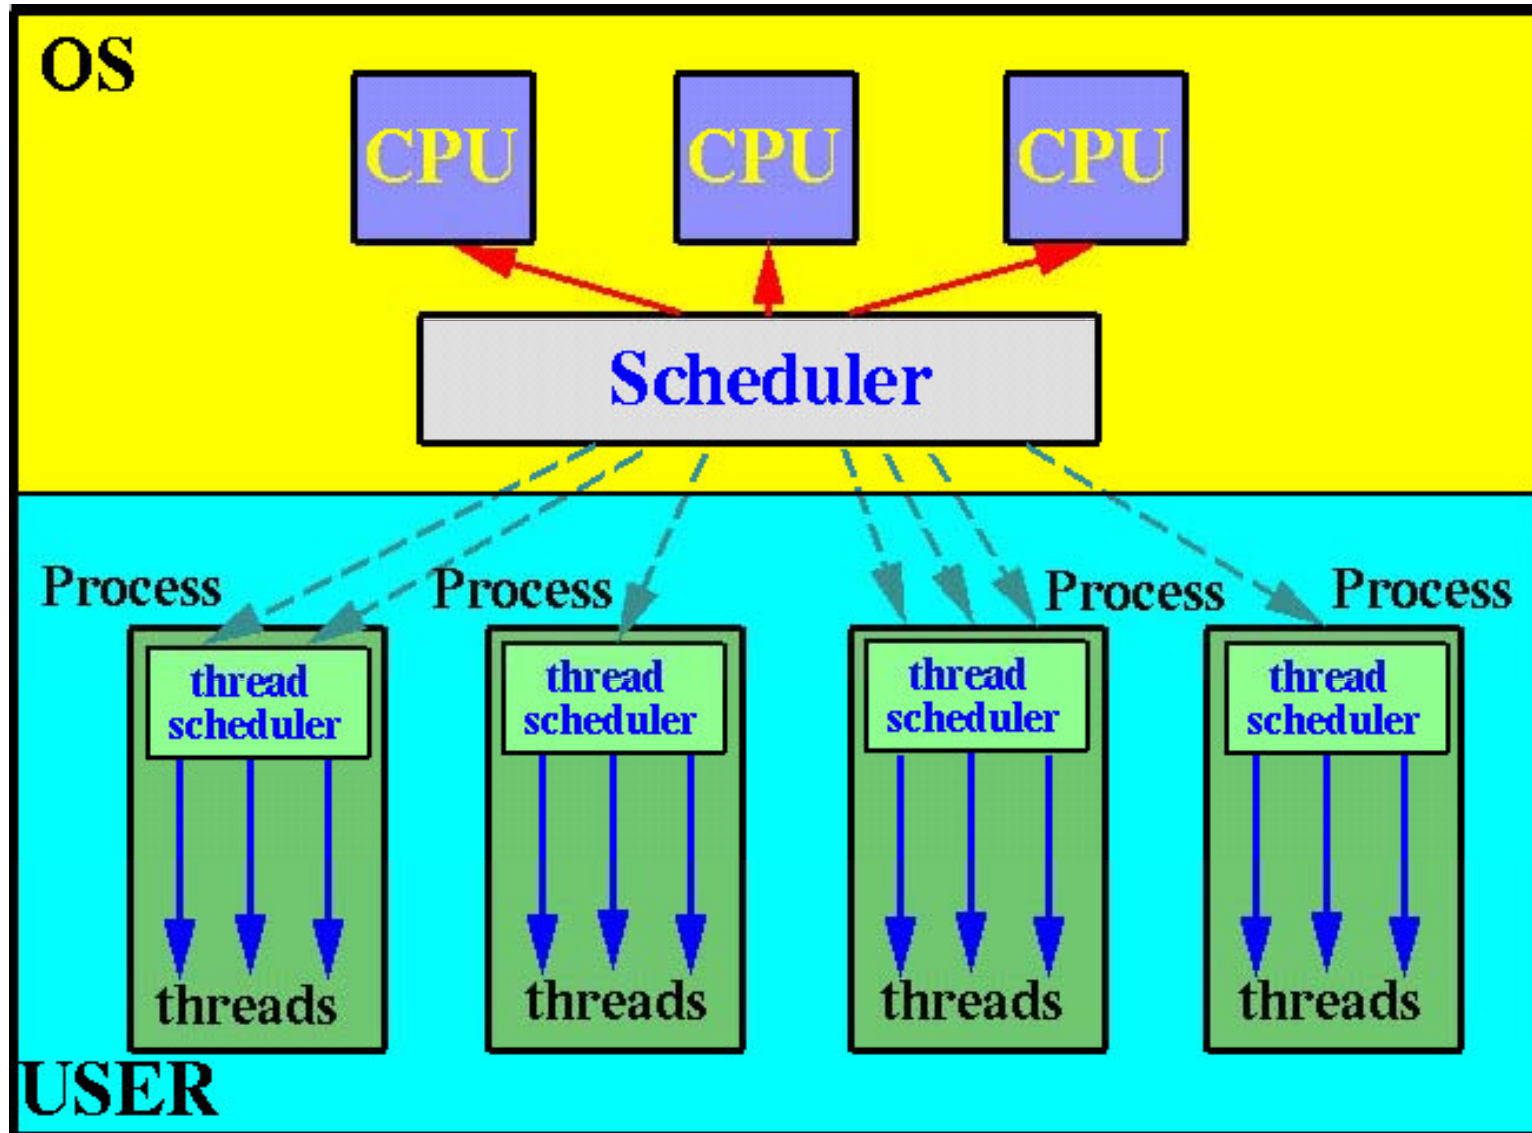
An Extreme Case: Traditional Unix



One-to-One Model: 2/2



Many-to-Many Model



Multicore Programming: 1/6

- ❑ With a single-core CPU, threads are scheduled by a scheduler and can only run one at a time.
- ❑ With a multicore CPU, multiple threads can run at the same time, one on each core.
- ❑ Therefore, system design becomes more complex than one may expect.
- ❑ Five issues have to be addressed properly: **dividing activities, balance, data splitting, data dependency, and testing and debugging.**


Multicore Programming: 2/6

□ Dividing Activities:

Since each thread can run on a core, one must study the problem in hand so that program activities can be divided and run concurrently.

□ Matrix multiplication is a good example.

□ Unfortunately, some problems are *inherently sequential* (e.g., DFS).




$$C_{i,j} = \sum_{k=1}^n A_{i,k} \times B_{k,j}$$

Multicore Programming: 3/6

- ❑ **Balance:** Make sure each thread has *equal* contribution, if possible, to the whole computation.
- ❑ If an insignificant thread runs frequently, occupying a core, other more useful threads would have less chance to run.

Multicore Programming: 4/6

- ❑ **Data Splitting:** Data may also be split into different sections so that they can be processed separately.
- ❑ Matrix multiplication is a good example.
- ❑ Quicksort is another. After partitioning, the two sections can be sorted separately.



Multicore Programming: 5/6

- ❑ **Data Dependency:** Watch for data items that are used by different threads. For example, two threads may update a common counter `count` with `count++`.
- ❑ Should this happen, unexpected results may occur. As a result, the execution of threads has to be synchronized so that only one thread can update `count` at any time.
- ❑ This is the most difficult part in threaded programming. See **Chapter 6**.

Multicore Programming: 6/6

- ❑ **Testing and Debugging:** The behavior of a threaded program is *dynamic*. A bug that appears in this test run may not occur in the next. Some bugs may never surface throughout the life-span of a threaded program, or may appear at an unexpected time.
- ❑ Some debugging issues (*e.g.*, race condition – updating a shared resource at the same time, and system deadlock) do not have efficient solutions.
- ❑ Thus, testing and debugging is an art, and requires a careful design and planning.

Thread Issues

- **How Does a Thread Fork?**
- **Thread Cancellation**
- **Signal Handling**
- **Thread-Specific Data**
- **What Is Thread-Safe?**

How Does a Thread Fork?

- ❑ If a thread forks, does the new process:
 - ❖ **duplicate** all threads?
 - ❖ contain only the forking thread (*i.e.*, single-threaded)?
- ❑ Some systems have **two** `fork` system calls, one for each case.

Thread Cancellation: 1/2

- ❑ ***Thread cancellation*** means terminating a thread **before** its completion. The thread that is to be cancelled is the ***target thread***.
- ❑ There are two types:
 - ❖ ***Asynchronous Cancellation***: the target thread terminates immediately.
 - ❖ ***Deferred Cancellation***: The target thread can periodically check if it should terminate, allowing the target thread an opportunity to terminate itself in an orderly way. The point a thread can terminate itself is a ***cancellation point***.

Thread Cancellation: 2/2

- ❑ **Problem:** With asynchronous cancellation, if the target thread owns some system-wide resources, the system may not be able to reclaim all resources owned by the target thread.
- ❑ With deferred cancellation, the target thread determines the time to terminate itself. Reclaiming resources is not a problem.
- ❑ Most systems implement asynchronous cancellation for processes (*e.g.*, system call `kill`) and threads.
- ❑ POSIX Threads (*i.e.*, Pthreads) supports deferred cancellation.

Signal Handling

- *Signals* is a way the OS uses to notify a process that some event has happened.
- Once a signal occurs, who is going to handle it? The process, or one of the threads?
- This is a very complex issue and will be discussed later in this course.

Thread-Specific Data/Thread-Safe

- ❑ Data that a thread needs for its own operation are *thread-specific*.
- ❑ Poor support for thread-specific data could cause problems. For example, while threads have their own stacks, they share the heap.
- ❑ What if two `malloc()`'s or `new`'s are executed at the same time requesting for memory from the heap? Or, two `printf`'s or `cout`'s are run simultaneously?
- ❑ If a library can be used by multiple threads properly, it is a *thread-safe* one.

Thread Pool

- ❑ While we know that managing threads are more efficient than managing processes, creating and terminating threads are still not free.
- ❑ After a process is created, one can immediately create a number of threads and have them waiting.
- ❑ When a new task occurs, one can wake up one of the waiting threads and assign it the work. After this thread completes the task, it goes back to wait.
- ❑ In this way, we save the number of thread creation and termination.
- ❑ These threads are said in a *thread pool*.

The End