

Part III

Storage Management

Chapter 10: File-System Interface

Files

- ❑ A file is a named collection of related information that is recorded on secondary storage.
- ❑ The operating systems maps this logical storage unit to the physical view of information storage.
- ❑ A file may have the following characteristics
 - ❖ File Attributes
 - ❖ File Operations
 - ❖ File Types
 - ❖ File Structures
 - ❖ Internal Files

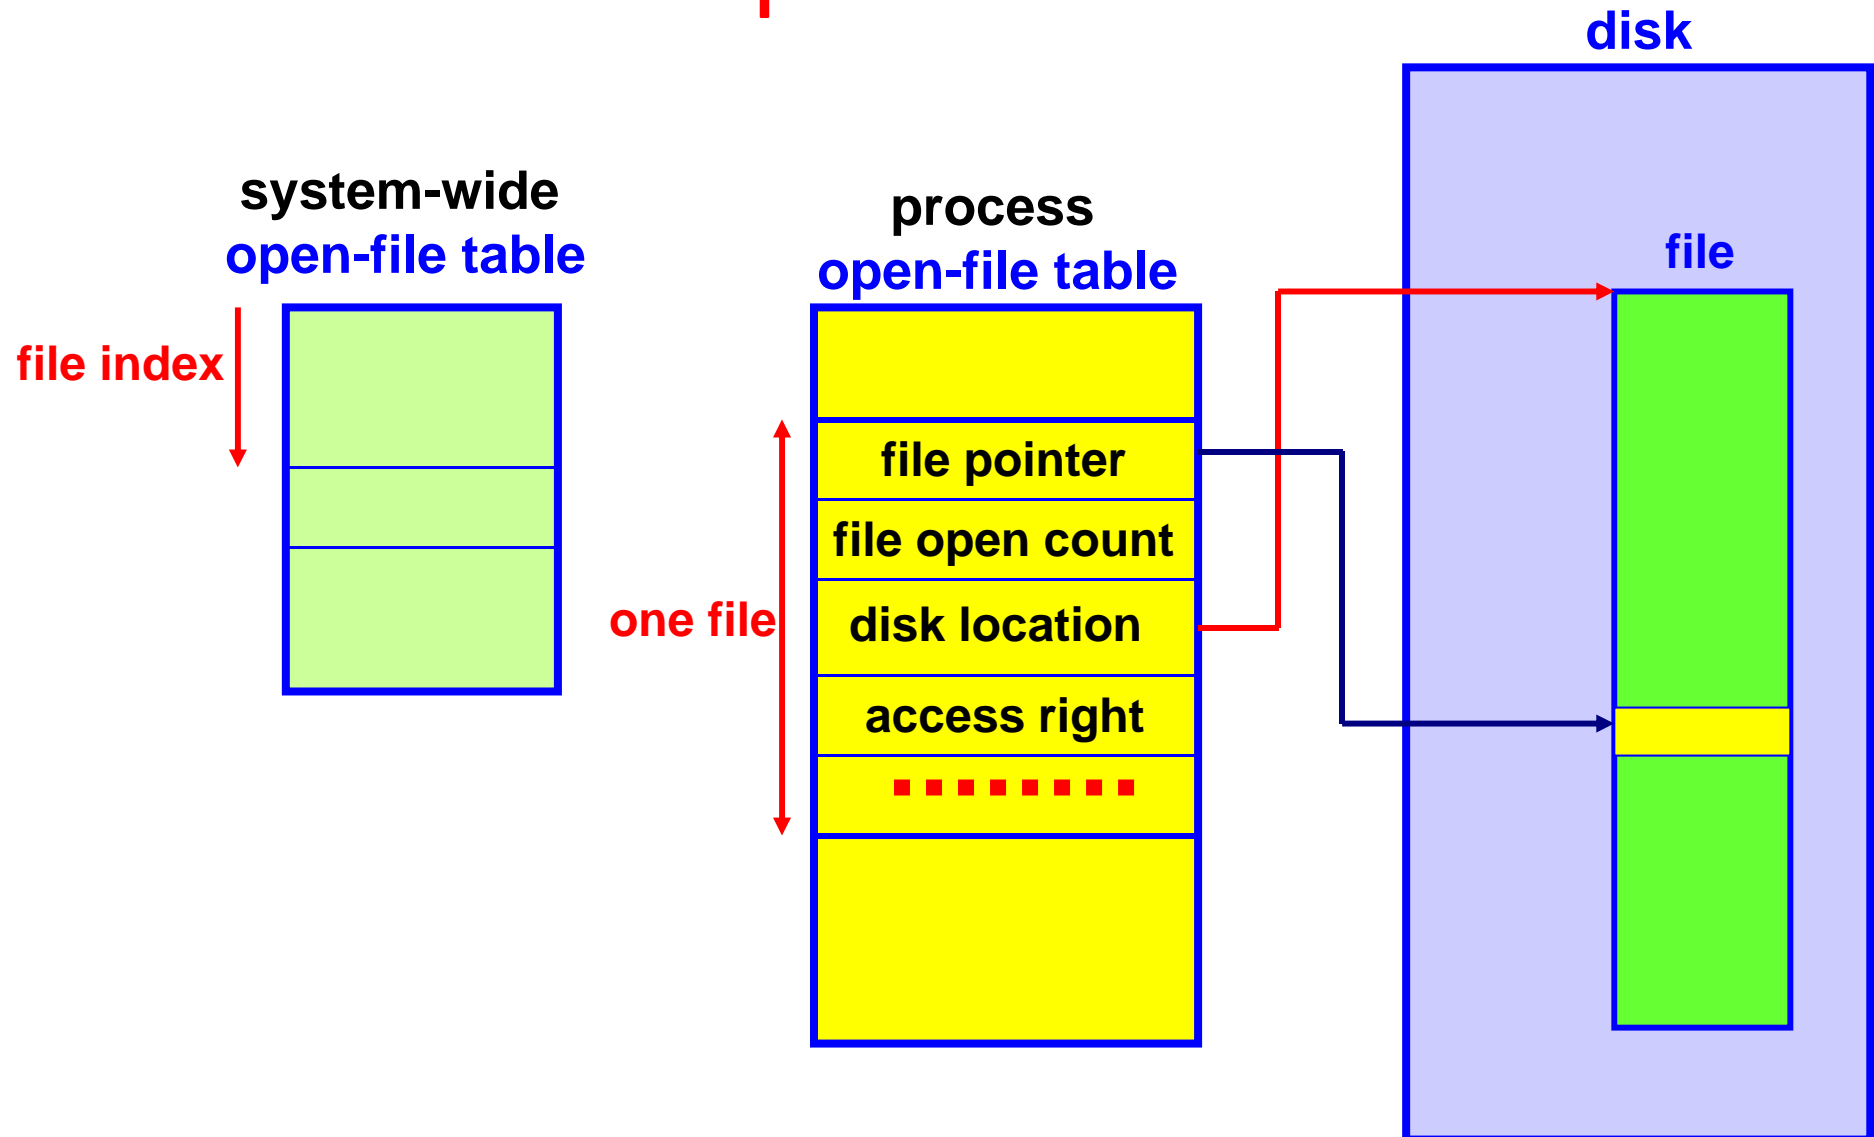
File Attributes

- ❑ **File Name:** The symbolic name is perhaps the only human readable file attribute.
- ❑ **Identifier:** A unique number assigned to each file for identification purpose.
- ❑ **File Type:** Some systems recognize various file types. Windows is a good example.
- ❑ **File Location:** A pointer to a device to find a file.
- ❑ **File Size:** The current size of a file, or the maximum allowed size.
- ❑ **File Protection:** This is for access-control.
- ❑ **File Date, Time, Owner, etc.**

File Operations: 1/2

- ☐ A file can be considered as an abstract data type that has data and accompanying operations.
- ☐ Creating a file
- ☐ Writing a file
- ☐ Reading a file
- ☐ Repositioning within a file
- ☐ Deleting a file
- ☐ Truncating a file
- ☐ Other operations (*e.g.*, appending a file, renaming a file)

File Operations: 2/2



File Structure


- ☐ Some systems support specific file types that have special file structures.
- ☐ For example, files that contain binary executables.
- ☐ An operating system becomes more complex when more file types (*i.e.*, file structures) are supported.
- ☐ In general, the number of supported file types is kept to minimum.

File Access Methods

- ❑ Access method: how a file be used.
- ❑ There are three popular ones:
 - ❖ Sequential access method for sequential files
 - ❖ Direct access method for direct files
 - ❖ Indexed access method for indexed files.

Sequential Access Method

- ❑ With the sequential access method, a file is processed in order, one record after the other.
- ❑ If p is the file pointer, the next record to be accessed is either $p+1$ (forward) or $p-1$ (i.e., backward).



Direct Access Method

- ❑ A file is made up of **fixed-length logical records**.
- ❑ The direct access method uses a **record number** to identify each record. For example, *read rec 0*, *write rec 100*, *seek rec 75*, etc.
- ❑ Some systems may use a **key field** to access a record (e.g., *read rec “Age=24”* or *write rec “Name=Dow”*). This is usually achieved with **hashing**.
- ❑ Since records can be accessed in random order, direct access is also referred to as **random access**.
- ❑ Direct access method can simulate sequential access.

Indexed Access Method

- ❑ With the **indexed access** method, a file is sorted in **ascending** order based on a number of keys.
- ❑ Each disk **block** may contain a number of **fixed-length logical records**.
- ❑ An **index table** stores the keys of the first block in each block.
- ❑ We can search the index table to locate the block that contains the desired record. Then, search the block to find the desired record.
- ❑ **This is exactly a one-level B-, B+ or B* tree.**
- ❑ **Multi-level index access method is also possible.**

index table

last name	logical rec #
Adams	
Arthur	
Ashcroft	
Smith	

*index tables are stored
in physical memory
when file is open*

data file

Ashcroft, ...	Asher, ...	Atkins
Smith,	Sweeny, ...	Swell, ...

Directory Structure: 1/2

- ❑ A large volume disk may be partitioned into **partitions**, or **mini disks**, or **volumes**.
- ❑ Each partition contains information about files within it. This information is stored in entries of a **device directory** or **volume table of content** (VTOC).
- ❑ The device directory, or **directory** for short, stores the name, location, size, type, access method, etc of each file.
- ❑ Operations perform on directory: **search for a file, create a file, delete a file, rename a file, traverse the file system**, etc.

Directory Structure: 2/2

□ There are five commonly used directory structures:

❖ Single-Level Directory

❖ Two-Level Directory


❖ Tree-Structure Directories

❖ Acyclic-Graph Directories

❖ General Graph Directories

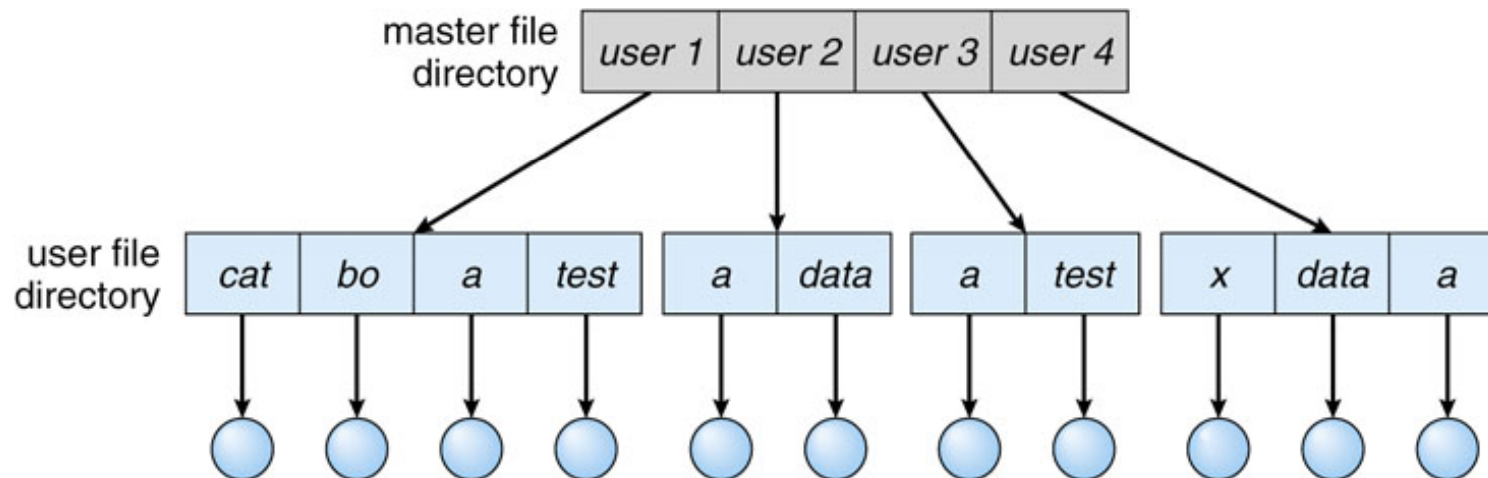
Single-Level Directory

- ❑ All files are contained in the same directory.
- ❑ It is difficult to maintain file name uniqueness.
- ❑ CP/M-80 and early version of MS-DOS use this directory structure.



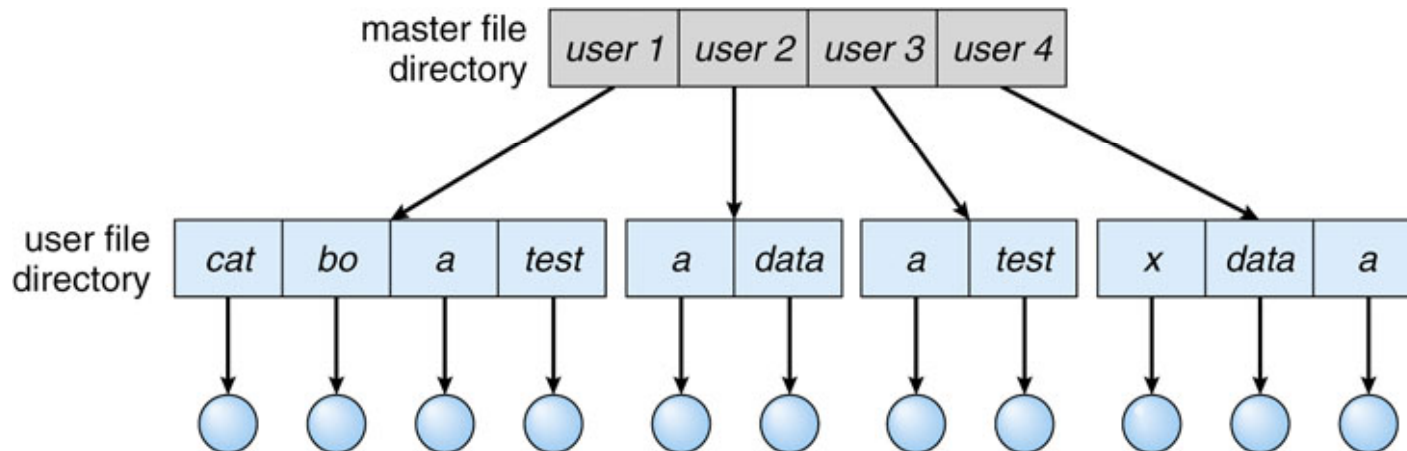
Two-Level Directory: 1/2

- ❑ This is an extension of the single-level directory for multi-user system.
- ❑ Each user has his/her user file directory. The system's master file directory is searched for the user directory when a user job starts.
- ❑ Early CP/M-80 multi-user systems use this structure.



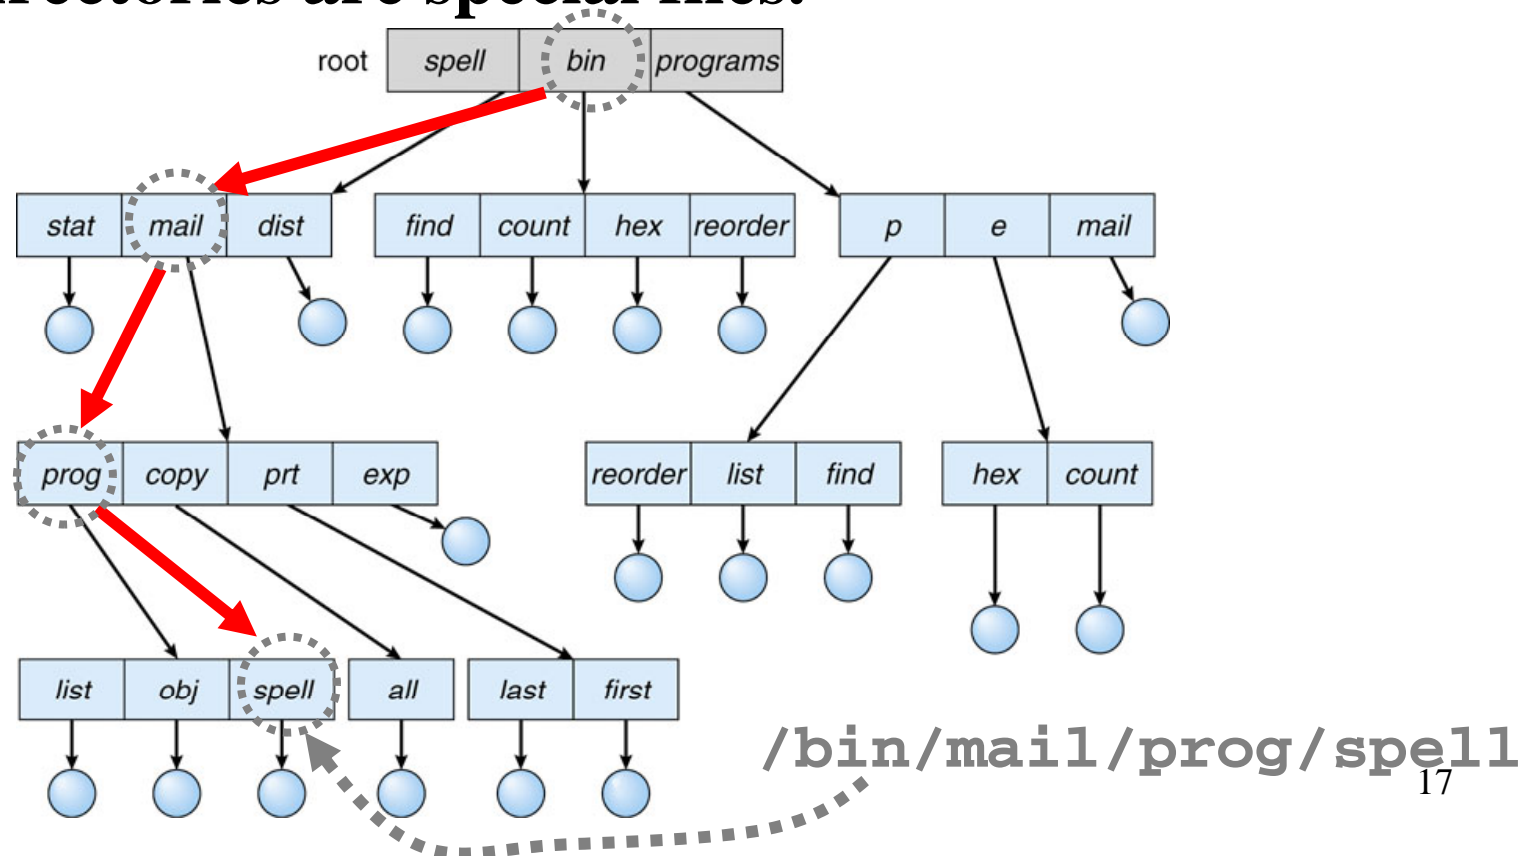
Two-Level Directory: 2/2

- ❑ To locate a file, path name is used. For example, `/user2/a` is the file `a` of `user 2`.
- ❑ Different systems use different path names. For example, under MS-DOS it is `C:\user2\a`.
- ❑ The directory of a special user, say `user 0`, may contain all system files.



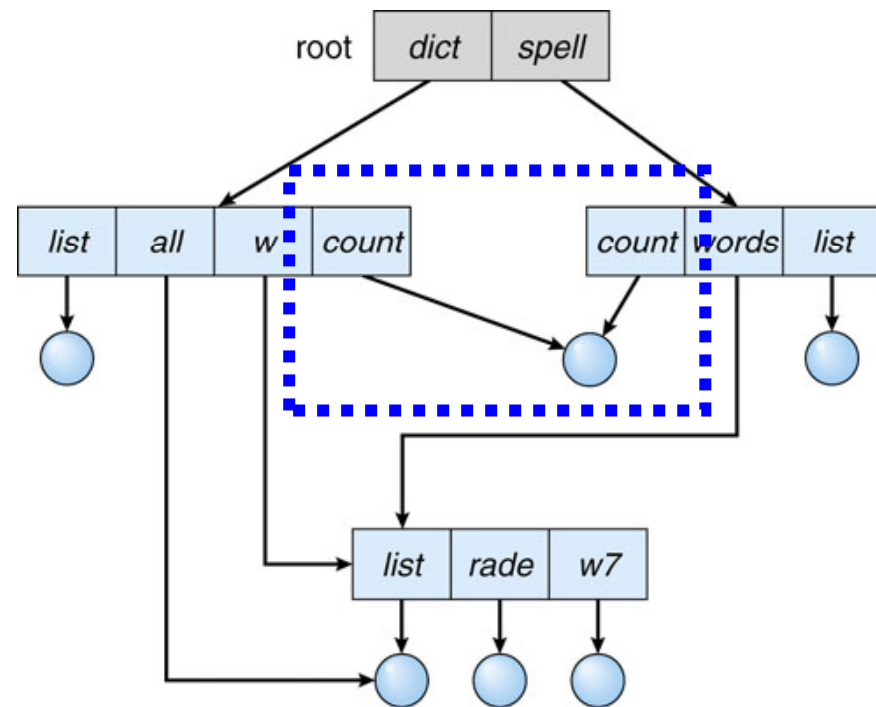
Tree-Structured Directory

- ❑ Each directory or subdirectory contains files and subdirectories, and forms a **tree**.
- ❑ Directories are special files.



Acyclic-Graph Directory: 1/2

- ❑ This type of directories allows a file/directory to be **shared** by multiple directories.
- ❑ This is different from two copies of the same file or directory.
- ❑ An acyclic-graph directory is more flexible than a simple tree structure. However, it is more complex.



file count is shared by directories
dict and spell

Acyclic-Graph Directory: 2/2

- ❑ Since a file have multiple absolute path names, how do we calculate file system statistics or do backup?


Would the same file be counted multiple times?

- ❑ How do we delete a file?

- ❖ If sharing is implemented with **symbolic links**, we only delete the link if we have a **list of links to the file**. The file is removed when the list is empty.
- ❖ Or, we **remove the file** and keep the links. When the file is accessed again, a message is given and the link is removed.
- ❖ Or, we can maintain a **reference count** for each shared file. The file is removed when the count is zero.

General Graph Directory: 1/2

- ❑ It is easy to traverse the directories of a tree or an acyclic directory system.
- ❑ However, if links are added arbitrarily, the directory graph becomes arbitrary and may contain **cycles**.
- ❑ *How do we search for a file?*



General Graph Directory: 2/2

- ❑ **How do we delete a file?** We can use **reference count**!
 - ❖ In a cycle, due to **self-reference**, the reference count may be non-zero even when it is no longer possible to refer to a file or directory.
 - ❖ Thus, **garbage collection** may be needed. A garbage collector traverses the directory and marks files and directories that can be accessed.
 - ❖ A second round removes those inaccessible items.
- ❑ To avoid this time-consuming task, a system can check if a cycle may occur when a link is made. How? You should know!

File Sharing

- ☐ When a file is shared by multiple users, how can we ensure its consistency?
- ☐ If multiple users are writing to the file, should all of the writers be allowed to write?
- ☐ Or, should the operating system protect the user actions from each other?
- ☐ This is the **file consistency semantics**.

File Consistency Semantics

- ❑ **Consistency semantics** is a characterization of the system that specifies the semantics of multiple users accessing a *shared* file *simultaneously*.
- ❑ Consistency semantics is an important criterion for evaluating any file system that supports file sharing.
- ❑ There are three commonly used semantics
 - ❖ **Unix semantics**
 - ❖ **Session Semantics**
 - ❖ **Immutable-Shared-Files Semantics**
- ❑ A **file session** consists all file accesses between `open()` and `close()`.

Unix Semantics

- ❑ Writes to an open file by a user are **visible immediately** to other users who have the file open at the **same** time.
- ❑ **All users share the file pointer.** Thus, advancing the file pointer by one user **affects all** sharing users.
- ❑ A file has a **single** image that interleaves all accesses, regardless of their origin.
- ❑ File access contention may cause delays.

Session Semantics

- ❑ Writes to an open file by a user are **not** visible immediately to other users who have the same file open simultaneously.
- ❑ Once a file is **closed**, the changes made to it are visible only **in sessions started later**.
- ❑ **Already-open instances of the file are not affected by these changes.**
 - ❖ A file may be associated **temporarily** with several and possible different images at the same time.
 - ❖ Multiple users are allowed to perform both read and write concurrently on **their** image of the file without delay.
- ❑ The Andrew File System (AFS) uses this semantics.

Immutable-Shared-Files Semantics

- ❑ Once a file is declared as shared by its creator, it **cannot** be modified.
- ❑ An immutable file has two important properties:
 - ❖ Its name may not be **used**
 - ❖ Its content may not be **altered**
- ❑ Thus, the name of an immutable file indicates that the contents of the file is **fixed** – a constant rather than a variable.
- ❑ The implementation of these semantics in a distributed system is simple, since sharing is disciplined (*i.e.*, **read-only**).

File Protection

- ❑ We can keep files safe from physical damage (*i.e.*, **reliability**) and improper access (*i.e.*, **protection**).
- ❑ **Reliability** is generally provided by backup.
- ❑ The need for **file protection** is a direct result of the ability to access files.
- ❑ Access control may be a **complete protection** by denying access. Or, the access may be **controlled**.

File Protection: Types of Access

- ❑ Access control may be implemented by limiting the types of file access that can be made.
- ❑ The types of access may be
 - ❖ **Read**: read from the file
 - ❖ **Write**: write or rewrite the file
 - ❖ **Execute**: load the file into memory and execute it
 - ❖ **Append**: write new info at the end of a file
 - ❖ **Delete**: delete a file
 - ❖ **List**: list the name and attributes of the file

File Protection: Access Control: 1/4





- ❑ The most commonly used approach is to make the access dependent on the identity of the user.
- ❑ Each file and directory is associated with an **access matrix** specifying the user name and the types of permitted access.
- ❑ When a user makes a request to access a file or a directory, his/her identity is compared against the information stored in the access matrix.

File Protection: Access Control: 2/4

Access Matrix

	File 1	File 2	File 3	File 4	Account 1	Account 2
User A	Own R W		Own R W		Inquiry Credit	
User B	R	Own R W	W	R	Inquiry debit	Inquiry Credit
User C	R W	R		Own R W		Inquiry debit

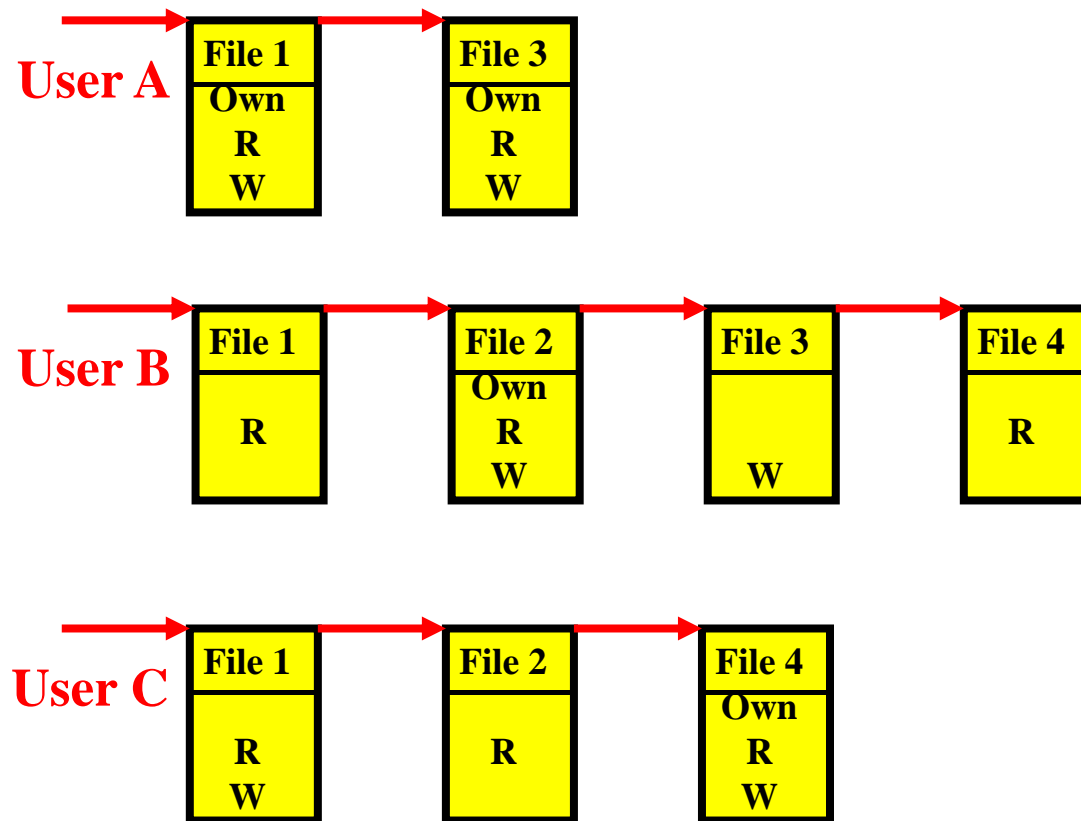
File Protection: Access Control: 3/4



Access-control Lists

- ☐ In practice, the access matrix is sparse.
- ☐ The matrix can be decomposed into **columns** (files), yielding **access-control lists (ACL)**
- ☐ However, this list can be very long!

File Protection: Access Control: 4/4



Capability Lists

- ❑ Decomposition by **rows** (users) yields **capability tickets**.
- ❑ Each user has a number of tickets for file/directory access.
- ❑ These tickets may be authorized to loan or be given to other users.
- ❑ All tickets may be held and managed by the OS for better protection.

The End